

Improving performance of C++ modules in Clang

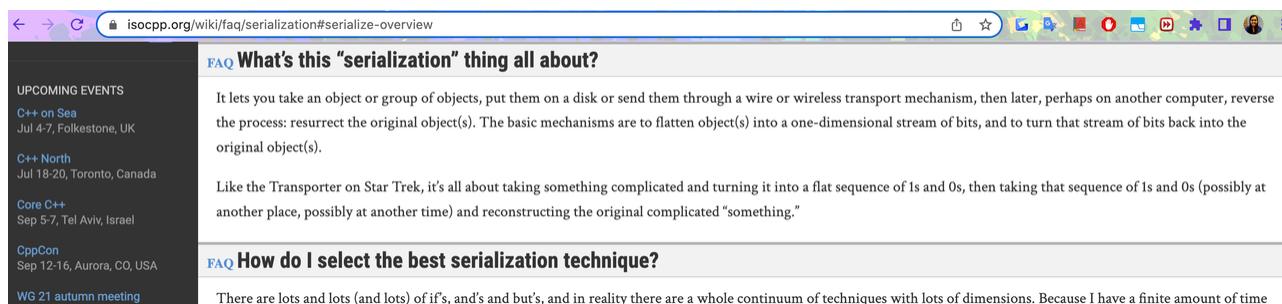
Problem Statement

The C++ modules technology aims to provide a scalable compilation model for the C++ language. The C++ Modules technology in Clang provides an io-efficient, on-disk representation capable to reduce build times and peak memory usage. The internal compiler state such as the **abstract syntax tree (AST) is stored on disk and lazily loaded on demand**. C++ Modules improve the memory footprint for interpreted C++ through the Cling C++ interpreter developed by CERN and the compiler research group at Princeton. The current implementation is pretty good at making most operations on demand.

However in a few cases, **we eagerly load pieces of the AST**, for example **at module import time** and upon **selecting a suitable template specialization**. When selecting the template specialization **we load all template specializations from the module files just to find out they are not suitable**. There is a patch that partially solves this issue by introducing a template argument hash and use it to look up the candidates without deserializing them. However, **the data structure it uses to store the hashes leads to quadratic search which is inefficient when the number of modules becomes sufficiently large**.

Serialization

Serialization is the process of writing or reading an object to or from a persistent storage medium such as a disk file.



The screenshot shows a web browser window with the URL isocpp.org/wiki/faq/serialization#serialize-overview. The page content is as follows:

UPCOMING EVENTS

- C++ on Sea
Jul 4-7, Folkestone, UK
- C++ North
Jul 18-20, Toronto, Canada
- Core C++
Sep 5-7, Tel Aviv, Israel
- CppCon
Sep 12-16, Aurora, CO, USA
- WG 21 autumn meeting

FAQ What's this "serialization" thing all about?

It lets you take an object or group of objects, put them on a disk or send them through a wire or wireless transport mechanism, then later, perhaps on another computer, reverse the process: resurrect the original object(s). The basic mechanisms are to flatten object(s) into a one-dimensional stream of bits, and to turn that stream of bits back into the original object(s).

Like the Transporter on Star Trek, it's all about taking something complicated and turning it into a flat sequence of 1s and 0s, then taking that sequence of 1s and 0s (possibly at another place, possibly at another time) and reconstructing the original complicated "something."

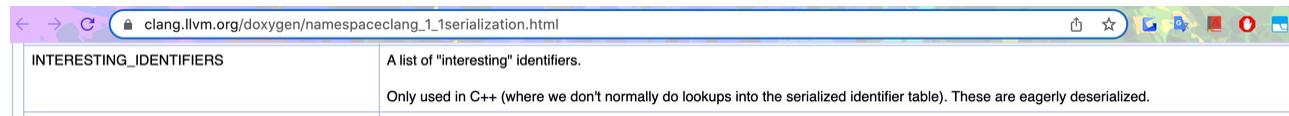
FAQ How do I select the best serialization technique?

There are lots and lots (and lots) of if's, and's and but's, and in reality there are a whole continuum of techniques with lots of dimensions. Because I have a finite amount of time

Deserialization

The byte stream, once created, also can be streamed across a communication link to a remote receiving end. The reverse of serialization is called deserialization, where the data in the byte stream is used to reconstruct it to its original object form.

Eager Deserialization



INTERESTING_IDENTIFIER

A list of "interesting" identifiers.

Only used in C++ (where we don't normally do lookups into the serialized identifier table). These are eagerly deserialized.

Example when eager deserialization cannot be avoided: until c++20 we could lazily deserialize the vtable information but due to `constexpr virtual` in c++20 we cannot anymore.

In Clang

Serialization and Deserialization

`clang/include/clang/Serialization/ASTDeserializationListener.h`

```

#include "clang/Basic/IdentifierTable.h"
#include "clang/Serialization/ASTBitCodes.h"

namespace clang {

class Decl;
class ASTReader;
class QualType;
class MacroDefinitionRecord;
class MacroInfo;
class Module;
class SourceLocation;

class ASTDeserializationListener {
public:
    virtual ~ASTDeserializationListener();

    /// The ASTReader was initialized.
    virtual void ReaderInitialized(ASTReader *Reader) { }

    /// An identifier was deserialized from the AST file.
    virtual void IdentifierRead(serialization::IdentID ID,
                                IdentifierInfo *II) { }

    /// A macro was read from the AST file.
    virtual void MacroRead(serialization::MacroID ID, MacroInfo *MI) { }

    /// A type was deserialized from the AST file. The ID here has the
    /// qualifier bits already removed, and T is guaranteed to be locally
    /// unqualified.
    virtual void TypeRead(serialization::TypeID Idx, QualType T) { }

    /// A decl was deserialized from the AST file.
    virtual void DeclRead(serialization::DeclID ID, const Decl *D) { }

    /// A selector was read from the AST file.
    virtual void SelectorRead(serialization::SelectorID iD, Selector Sel) {}

    /// A macro definition was read from the AST file.
    virtual void MacroDefinitionRead(serialization::PreprocessedEntityID,
                                     MacroDefinitionRecord *MD) {}

    /// A module definition was read from the AST file.
    virtual void ModuleRead(serialization::SubmoduleID ID, Module *Mod) {}

    /// A module import was read from the AST file.
    virtual void ModuleImportRead(serialization::SubmoduleID ID,
                                 SourceLocation ImportLoc) {}

};

}

```

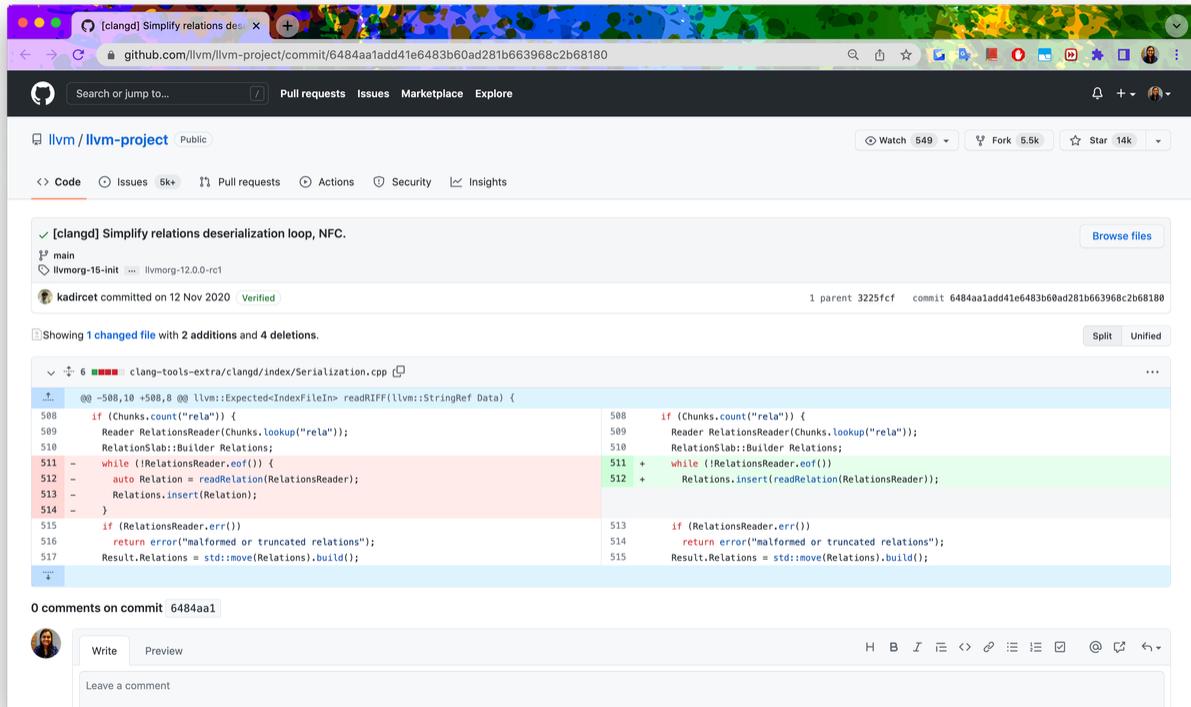
```

void ASTWriter::ModuleRead(serialization::SubmoduleID ID, Module *Mod) {
    assert(SubmoduleIDs.find(Mod) == SubmoduleIDs.end());
    SubmoduleIDs[Mod] = ID;
}

```

}

Simple code to understand deserialization



[clangd] Simplify relations deserialization loop, NFC.

main
llvmorg-15-init ... llvmorg-12.0.0-rct

kadiracet committed on 12 Nov 2020 · Verified

Showing 1 changed file with 2 additions and 4 deletions.

clang-tools-extra/clangd/index/Serialization.cpp

```
@@ -508,10 +508,8 @@ llvm::Expected<IndexFileIn> readRIFF(llvm::StringRef Data) {
508     if (Chunks.count("rela")) {
509         Reader RelationsReader(Chunks.lookup("rela"));
510         RelationSlab::Builder Relations;
511         while (!RelationsReader.eof()) {
512             auto Relation = readRelation(RelationsReader);
513             Relations.insert(Relation);
514         }
515         if (RelationsReader.err())
516             return error("malformed or truncated relations");
517         Result.Relations = std::move(Relations).build();
508     if (Chunks.count("rela")) {
509         Reader RelationsReader(Chunks.lookup("rela"));
510         RelationSlab::Builder Relations;
511         + while (!RelationsReader.eof())
512         +     Relations.insert(readRelation(RelationsReader));
513         if (RelationsReader.err())
514             return error("malformed or truncated relations");
515         Result.Relations = std::move(Relations).build();

```

1 parent 3225fcf commit 6484aa1add41e6483b60ad281b663968c2b68180

Split Unified

Comments on commit 6484aa1

Write Preview

Leave a comment

Eager Deserialization

Module import time

<https://github.com/llvm/llvm-project/commit/c52efa7d4011a48ea91b353f2cbc40a359d19571>

[modules] Don't eagerly deserialize so many ImportDecls. CodeGen basi...
...cally ignores ImportDecls imported from modules, so only eagerly deserialize the ones from a PCH / preamble.

llvm-svn: 245406

↳ main
↳ llvmlang-15-init ... 2020.06-alpha

zygoloid committed on 19 Aug 2015 1 parent 72be1c1 commit c52efa7d4011a48ea91b353f2cbc40a359d19571

Showing 2 changed files with 13 additions and 11 deletions. Split Unified

clang/lib/CodeGen/CodeGenModule.cpp

```
@@ -3363,11 +3363,8 @@ void CodeGenModule::EmitTopLevelDecl(Decl *D) {  
    auto *Import = cast<ImportDecl>(D);  
    // Ignore import declarations that come from imported modules.  
-   if (clang::Module *Owner = Import->getImportedOwningModule()) {  
-       if (getLangOpts().CurrentModule.empty() ||  
-           Owner->getTopLevelModule()->Name == getLangOpts().CurrentModule)  
-           break;  
-   }  
+   if (Import->getImportedOwningModule())  
+       break;  
  
    if (CGDebugInfo *DI = getModuleDebugInfo())  
        DI->EmitImportDecl(*Import);  
}
```

clang/lib/Serialization/ASTWriterDecl.cpp

```
@@ -1994,14 +1994,19 @@ void ASTWriter::WriteDeclAbbrevs() {  
    /// clients to use a separate API call to "realize" the decl. This should be  
    1994    /// clients to use a separate API call to "realize" the decl. This should be
```

Upon selecting a suitable template specialization

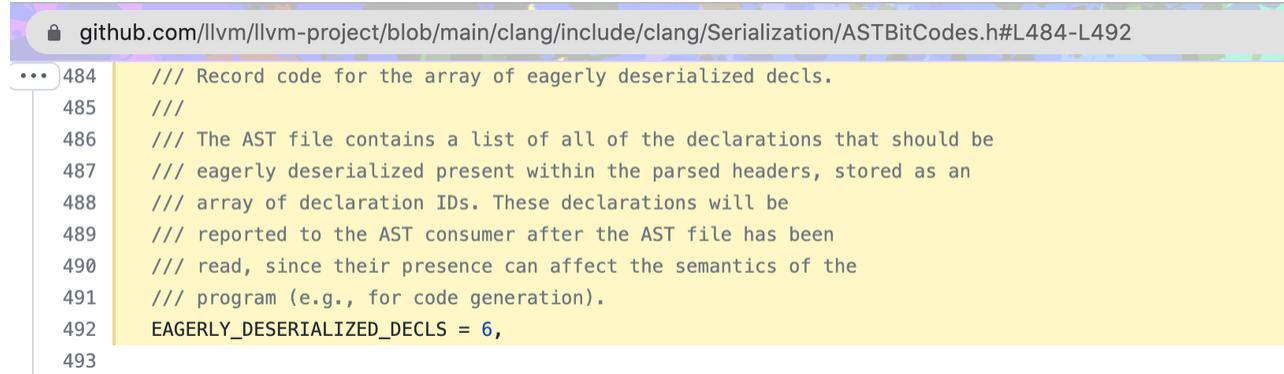
When selecting the template specialization we load all template specializations from the module files just to find out they are not suitable.

With lazy deserialization, builtins are loaded on-demand, and unused builtins are never loaded into the Isolate. Lazy deserialization comes with memory savings.

Existing (using print statements)

<https://github.com/llvm/llvm-project/blob/main/clang/include/clang/Serialization/ASTBitCodes.h#L484-L492>

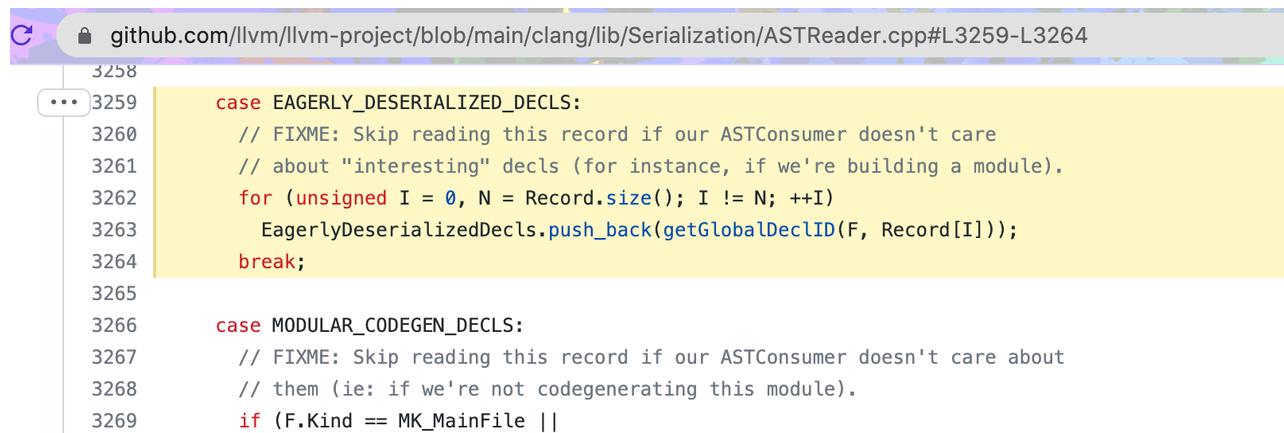
EAGERLY_DESERIALIZED_DECLS



github.com/llvm/llvm-project/blob/main/clang/include/clang/Serialization/ASTBitCodes.h#L484-L492

```
... 484     /// Record code for the array of eagerly serialized decls.  
485     ///  
486     /// The AST file contains a list of all of the declarations that should be  
487     /// eagerly serialized present within the parsed headers, stored as an  
488     /// array of declaration IDs. These declarations will be  
489     /// reported to the AST consumer after the AST file has been  
490     /// read, since their presence can affect the semantics of the  
491     /// program (e.g., for code generation).  
492 EAGERLY_DESERIALIZED_DECLS = 6,  
493
```

1. <https://github.com/llvm/llvm-project/blob/main/clang/lib/Serialization/ASTReader.cpp#L3259-L3264>



github.com/llvm/llvm-project/blob/main/clang/lib/Serialization/ASTReader.cpp#L3259-L3264

```
3258  
... 3259     case EAGERLY_DESERIALIZED_DECLS:  
3260         // FIXME: Skip reading this record if our ASTConsumer doesn't care  
3261         // about "interesting" decls (for instance, if we're building a module).  
3262         for (unsigned I = 0, N = Record.size(); I != N; ++I)  
3263             EagerlySerializedDecls.push_back(getGlobalDeclID(F, Record[I]));  
3264         break;  
3265  
3266     case MODULAR_CODEGEN_DECLS:  
3267         // FIXME: Skip reading this record if our ASTConsumer doesn't care about  
3268         // them (ie: if we're not codegenerating this module).  
3269         if (F.Kind == MK_MainFile ||
```

2. <https://github.com/llvm/llvm-project/blob/main/clang/lib/Serialization/ASTReader.cpp#L7486>

```
github.com/llvm/llvm-project/blob/main/clang/lib/Serialization/ASTReader.cpp#L7486

... 7486     if (!DeclsLoaded[Index]) {
7487         ReadDeclRecord(ID);
7488         if (DeserializationListener)
7489             DeserializationListener->DeclRead(ID, DeclsLoaded[Index]);
7490     }
7491
7492     return DeclsLoaded[Index];
7493 }
7494
```

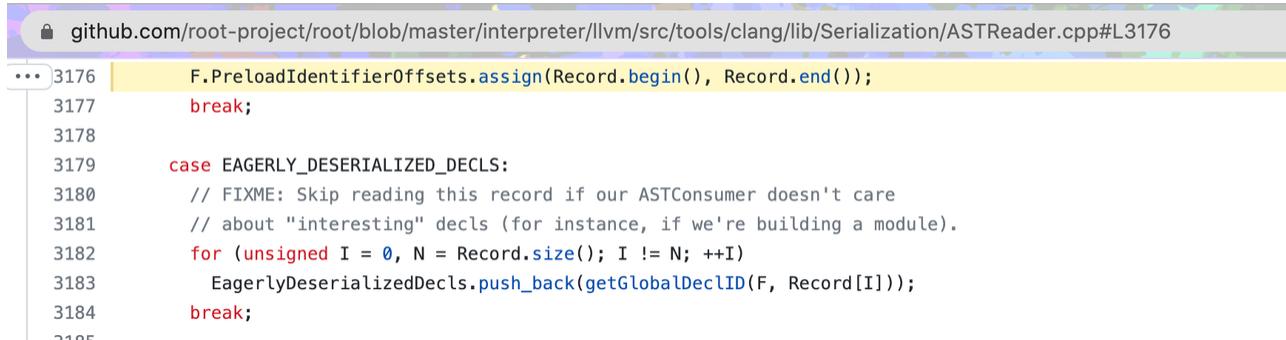
3. <https://github.com/llvm/llvm-project/blob/main/clang/lib/Serialization/ASTReader.cpp#L1573-L1604>

Preallocated source locations for modules which are not loaded. There was some plan to reduce this but didn't go anywhere.

```
github.com/llvm/llvm-project/blob/main/clang/lib/Serialization/ASTReader.cpp#L1573-L1604

... 1573     case SM_SLOC_BUFFER_ENTRY: {
1574         const char *Name = Blob.data();
1575         unsigned Offset = Record[0];
1576         SrcMgr::CharacteristicKind
1577             FileCharacter = (SrcMgr::CharacteristicKind)Record[2];
1578         SourceLocation IncludeLoc = ReadSourceLocation(*F, Record[1]);
1579         if (IncludeLoc.isInvalid() && F->isModule()) {
1580             IncludeLoc = getImportLocation(F);
1581         }
1582
1583         auto Buffer = ReadBuffer(SLocEntryCursor, Name);
1584         if (!Buffer)
1585             return true;
1586         SourceMgr.createFileID(std::move(Buffer), FileCharacter, ID,
1587                               BaseOffset + Offset, IncludeLoc);
1588         break;
1589     }
1590
1591     case SM_SLOC_EXPANSION_ENTRY: {
1592         LocSeq::State Seq;
1593         SourceLocation SpellingLoc = ReadSourceLocation(*F, Record[1], Seq);
1594         SourceLocation ExpansionBegin = ReadSourceLocation(*F, Record[2], Seq);
1595         SourceLocation ExpansionEnd = ReadSourceLocation(*F, Record[3], Seq);
1596         SourceMgr.createExpansionLoc(SpellingLoc, ExpansionBegin, ExpansionEnd,
1597                                     Record[5], Record[4], ID,
1598                                     BaseOffset + Record[0]);
1599         break;
1600     }
1601
1602     return false;
1603 }
1604
```

4. Another preloading: <https://github.com/root-project/root/blob/master/interpreter/llvm/src/tools/clang/lib/Serialization/ASTReader.cpp#L3176>



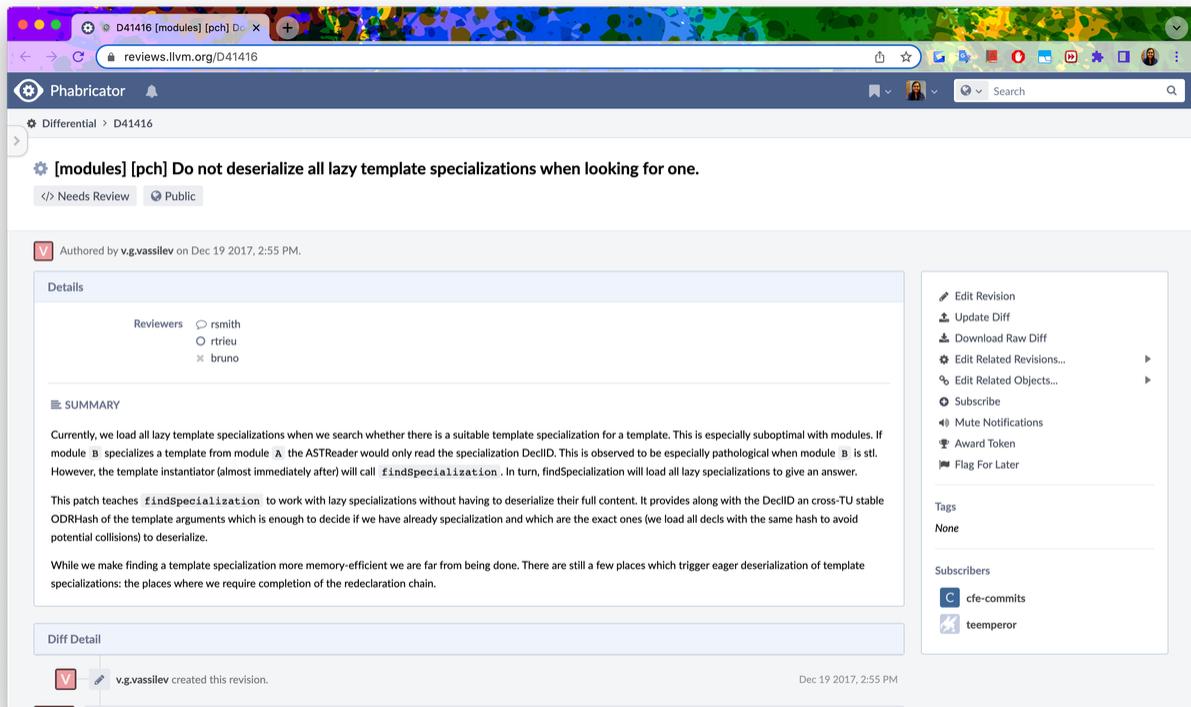
```

3176     F.PreloadIdentifierOffsets.assign(Record.begin(), Record.end());
3177     break;
3178
3179     case EAGERLY_DESERIALIZED_DECLS:
3180         // FIXME: Skip reading this record if our ASTConsumer doesn't care
3181         // about "interesting" decls (for instance, if we're building a module).
3182         for (unsigned I = 0, N = Record.size(); I != N; ++I)
3183             EagerlyDeserializedDecls.push_back(getGlobalDeclID(F, Record[I]));
3184         break;
3185

```

Previous work

<https://reviews.llvm.org/D41416>



The screenshot shows a Phabricator Differential interface for a patch titled "[modules] [pch] Do not deserialize all lazy template specializations when looking for one." The patch was authored by v.g.vassilev on Dec 19 2017, 2:55 PM. It has two reviewers listed: rsmith, rtrieu, and bruno. The summary explains that the patch teaches `findSpecialization` to work with lazy specializations without having to deserialize their full content. It provides along with the DeclID an cross-TU stable ODR-Hash of the template arguments which is enough to decide if we have already specialization and which are the exact ones (we load all decls with the same hash to avoid potential collisions) to deserialize.

The patch also notes that while making finding a template specialization more memory-efficient, there are still a few places which trigger eager deserialization of template specializations: the places where we require completion of the redeclaration chain.

The right sidebar contains various actions such as Edit Revision, Update Diff, Download Raw Diff, Edit Related Revisions..., Edit Related Objects..., Subscribe, Mute Notifications, Award Token, and Flag For Later. It also shows that there are no tags and two subscribers: cfe-commits and teemperor.

Partially solves this issue by introducing a template argument hash and use it to look up the candidates without deserializing them.

This way we managed to catch a few collisions in the ODRHash logic.

Check if we have already specialization and which are the exact ones (we load all decls with the same hash to avoid potential collisions) to deserialize.

Improvement/Optimization: the data structure it uses to store the hashes leads to quadratic search which is inefficient when the number of modules becomes sufficiently large.

Roadmap

Investigate and resolve eager deserialization where possible

1. Use the internal clang AST counters to file what is eagerly deserialize.
2. Add `printf` in `ASTReader::ReadDecl` and load a bunch of modules without using them. This ideally should be a nop. If that's not the case it has to be debugged and investigated further.

Rework the patch to use on-disk hash tables to avoid the quadratic search complexity

1. Move to using an on-disk hash table for template specialization lookup, at least for templates with large numbers of specializations
2. Currently when we hash a tag type the visitor calls `ODRHash::AddDecl` which mostly relies on the decl name give distinct hash value. The types coming from template specializations have very similar properties (including decl names). For those we need to provide more information in order to disambiguate them. This patch adds the template arguments for the template specialization decl corresponding to its type. We manage to reduce further the amount of deserializations from 1117 down to 451.

v.g.vassilev updated this revision to Diff 128802.

Reduce further the deserializations from 451 to 449 by providing a more complete implementation of `ODRHash::AddTemplateArgument`.

Kudos @rtrieu !

3. Stats:

- o types read is down from 30% to 17%
- o declarations read is down from 34% to 23%
- o number of ClassTemplateSpecializations read has decreased by 30%,
- o number of CXXRecordDecls read is down 25%
- o total ASTContext memory usage is down by 12%

4. calculate hash

```
104
705     unsigned TemplateArgumentList::ComputeODRHash(ArrayRef<TemplateArgument> Args) {
706         ODRHash Hasher;
707         for (TemplateArgument TA : Args)
708             Hasher.AddTemplateArgument(TA);
709
710         return Hasher.CalculateHash();
711     }
712
713 }
```

FunctionTemplateSpecializationInfo ↴

5. Add template argument

```
// If this was a specialization we should take into account its template
// arguments. This helps to reduce collisions coming when visiting template
// specialization types (eg. when processing type template arguments).
ArrayRef<TemplateArgument> Args;
if (auto *CTSD = dyn_cast<ClassTemplateSpecializationDecl>(D))
    Args = CTSD->getTemplateArgs().asArray();
else if (auto *VTS defense = dyn_cast<VarTemplateSpecializationDecl>(D))
    Args = VTS defense->getTemplateArgs().asArray();
else if (auto *FD = dyn_cast<FunctionDecl>(D))
    if (FD->getTemplateSpecializationArgs())
        Args = FD->getTemplateSpecializationArgs()->asArray();

for (auto &TA : Args)
    AddTemplateArgument(TA);
```

6. ASTWriter.cpp

6996 auto *DC = cast<DeclContext>(DCDecl);
6997 SmallVector<Decl*, Bx::Decls;
6998 FindExternalLexicalDecls(
6999 DC, [&](Decl::Kind K) { return K == D->getKind(); }, Decls);
7000 }
7001 }
7002 }
7003
7004 if (auto *CTSD = dyn_cast<ClassTemplateSpecializationDecl>(D))
7005 CTSD->getSpecializedTemplate()->LoadLazySpecializations();
7006 if (auto *VTS defense: TSD = dyn_cast<VarTemplateSpecializationDecl>(D))
7007 VTS->getSpecializedTemplate()->LoadLazySpecializations();
7008 if (auto *FD = dyn_cast<FunctionDecl>(D)) {
7009 if (auto *Template = FD->getPrimaryTemplate())
7010 Template->LoadLazySpecializations();
7011 }
7012 }
7013
6996 auto *DC = cast<DeclContext>(DCDecl);
6997 SmallVector<Decl*, Bx::Decls;
6998 FindExternalLexicalDecls(
6999 DC, [&](Decl::Kind K) { return K == D->getKind(); }, Decls);
7000 }
7001 }
7002 }
7003
7004 RedeclarableTemplateDecl *Template = nullptr;
7005 ArrayRef<TemplateArgument> Args;
7006 if (auto *CTSD = dyn_cast<ClassTemplateSpecializationDecl>(D)) {
7007 Template = CTSD->getSpecializedTemplate();
7008 Args = CTSD->getTemplateArgs(); asArray();
7009 } else if (auto *VTS defense: TSD = dyn_cast<VarTemplateSpecializationDecl>(D)) {
7010 Template = VTS->getSpecializedTemplate();
7011 Args = VTS->getTemplateArgs(); asArray();
7012 } else if (auto *FD = dyn_cast<FunctionDecl>(D)) {
7013 if (auto *Tmplt = FD->getPrimaryTemplate()) {
7014 Template = Tmplt;
7015 }
7016 Args = FD->getTemplateSpecializationArgs(); asArray();
7017 }
7018 }
7019
7020 if (Template)
7021 Template->loadLazySpecializationsImpl(Args);
7022 }

rsmith

7. Added template specialisation info.

8. <https://github.com/root-project/root/blob/master/interpreter/llvm/src/tools/clang/lib/Serialization/ASTReader.cpp#L3134-L3145>

```
github.com/root-project/root/blob/master/interpreter/llvm/src/tools/clang/lib/Serialization/ASTReader.cpp#L3134-L3145
...
3134     case IDENTIFIER_TABLE:
3135         F.IdentifierTableData = Blob.data();
3136         if (Record[0]) {
3137             F.IdentifierLookupTable = ASTIdentifierLookupTable::Create(
3138                 (const unsigned char *)F.IdentifierTableData + Record[0],
3139                 (const unsigned char *)F.IdentifierTableData + sizeof(uint32_t),
3140                 (const unsigned char *)F.IdentifierTableData,
3141                 ASTIdentifierLookupTrait(*this, F));
3142
3143             PP.getIdentifierTable().setExternalIdentifierLookup(this);
3144         }
3145         break;
3146
```

Read a blob of identifiers from a module file and then put that blob into that table which is of type `llvm::OnDiskIterableChainedHashTable`.

Measure performance improvements

Size – du -sh *pcm

sort largest to smallest measure of file space amount recursively stored in directory

Memory Consumption – /usr/bin/time -v root.exe -I -b -q tutorials/hsimple.C

Compared against eager deserialization, reduce heap size.

Use the internal performance counters in clang - <https://godbolt.org/z/s61fxoYPs>

Internal performance counters:

*** AST Context Stats:

25662 types total.

5 Decayed types, 48 each (240 bytes)

133 ConstantArray types, 56 each (7448 bytes)

21 DependentSizedArray types, 64 each (1344 bytes)

19 IncompleteArray types, 40 each (760 bytes)

62 Built-in types, 24 each (1488 bytes)

103 Decltype types, 40 each (4120 bytes)

18 Auto types, 48 each (864 bytes)

969 DependentName types, 48 each (46512 bytes)

43 DependentTemplateSpecialization types, 48 each (2064 bytes)

736 Elaborated types, 48 each (35328 bytes)

6419 FunctionProto types, 40 each (256760 bytes)

645 InjectedClassName types, 40 each (25800 bytes)

```
76 MemberPointer types, 48 each (3648 bytes)
148 PackExpansion types, 40 each (5920 bytes)
98 Paren types, 40 each (3920 bytes)
1861 Pointer types, 40 each (74440 bytes)
1505 LValueReference types, 40 each (60200 bytes)
324 RValueReference types, 40 each (12960 bytes)
1015 SubstTemplateTypeParm types, 40 each (40600 bytes)
87 Enum types, 32 each (2784 bytes)
716 Record types, 32 each (22912 bytes)
6815 TemplateSpecialization types, 40 each (272600 bytes)
2935 TemplateTypeParm types, 40 each (117400 bytes)
32 TypeOfExpr types, 32 each (1024 bytes)
869 Typedef types, 32 each (27808 bytes)
1 UnaryTransform types, 48 each (48 bytes)
7 Using types, 40 each (280 bytes)
```

Total bytes = 1029272

```
31/518 implicit default constructors created
98/591 implicit copy constructors created
54/543 implicit move constructors created
34/595 implicit copy assignment operators created
7/543 implicit move assignment operators created
43/544 implicit destructors created
```

Number of memory regions: 513

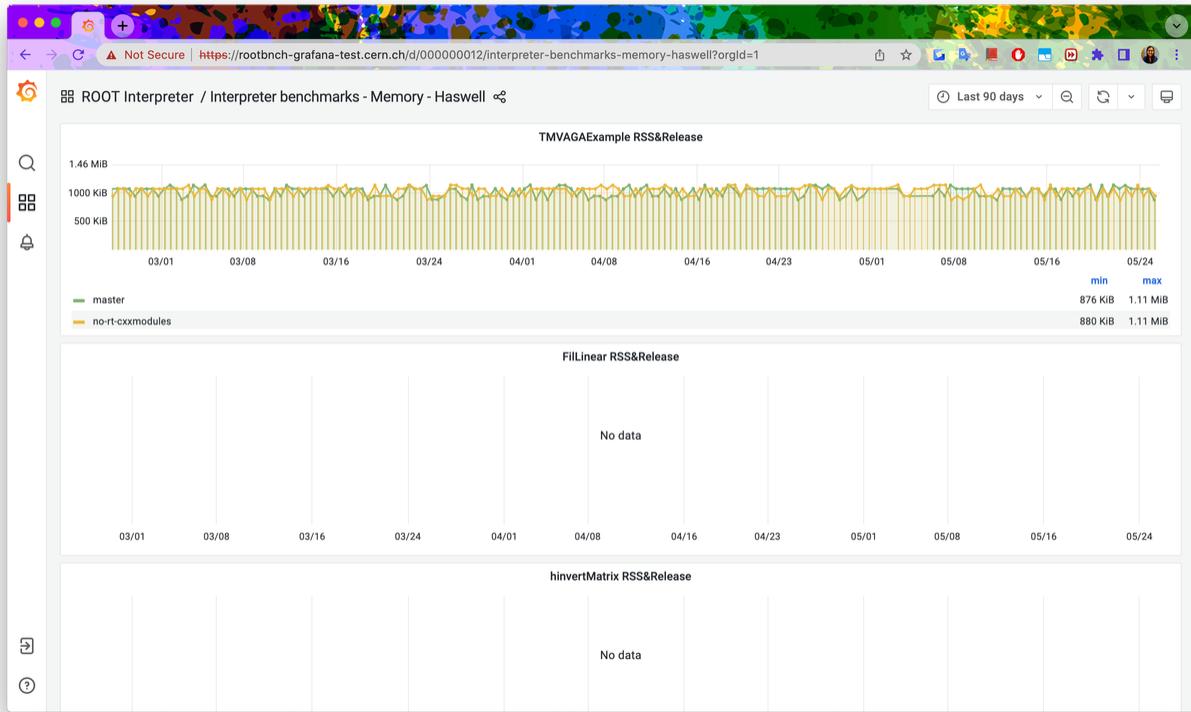
Bytes used: 7701107

Bytes allocated: 7929856

Bytes wasted: 228749 (includes alignment, etc)

Reduced memory consumption – ask Google to run the reimplementations of D41416 on their builds

<https://rootbnch-grafana-test.cern.ch/dashboards>



Build ROOT with `-Druntime_cxxmodules=On` on Windows

How to model the partial template specializations

Allows customizing class and variable templates for a given category of template arguments.

Examples of partial specializations in the standard library include `std::unique_ptr`, which has a partial specialization for array types.

example: from https://en.cppreference.com/w/cpp/language/partial_specialization

When a class or variable template is instantiated, and there are partial specializations available, the compiler has to decide if the primary template is going to be used or one of its partial specializations.

1) If only one specialization matches the template arguments, that specialization is used

2) If more than one specialization matches, partial order rules are used to determine which specialization is more specialized. The most specialized specialization is used, if it is unique (if it is not unique, the program cannot be compiled)

3) If no specializations match, the primary template is used

- the first function template has the same template parameters as the first partial specialization and has just one function parameter, whose type is a class template specialization with all the template arguments from the first partial specialization
- the second function template has the same template parameters as the second partial specialization and has just one function parameter whose type is a class template specialization with all the template arguments from the second partial specialization.

The function templates are then ranked as if for [function template overloading](#).