

cppyy

CaaS monthly meeting – 09/02/21

Wim Lavrijsen

cppyy: Yet another Python – C++ binder?!

- Yes, but it has its niche: *bindings are runtime*
 - Python is all runtime, so runtime is more natural
 - C++-side runtime-ness is provided by Cling
- Very complete feature-set (not just “C with classes”)
- Good performance on CPython; great with PyPy*

pip: <https://pypi.org/project/cppyy/>

conda: <https://anaconda.org/conda-forge/cppyy>

git: <https://github.com/wlav/cppyy>

docs: <https://cppyy.readthedocs.io/en/latest/>

For HEP users: *cppyy in ROOT is an old fork. It won't run all the examples here, doesn't work with PyPy, and has worse performance.*

(*) PyPy support lags CPython



Examples of Runtime Behavior

Runtime Template Instantiations

- **Cling instantiates templates at runtime**
 - No pre-instantiation/compilation necessary
 - Prevent duplication of standard classes (e.g. STL)
 - No combinatorial explosion (esp. with numeric types)
 - Support for templates of user classes

Runtime Template Instantiations

```
struct MyClass {  
    MyClass(int i) : fData(i) {}  
    virtual ~MyClass() {}  
    virtual int add(int i) {  
        return fData + i;  
    }  
    int fData;  
};
```

```
>>> import cppyy.gbl as CC  
>>> v = \  
...     CC.std.vector[CC.MyClass] ()  
...  
>>> for i in range(10):  
...     v.emplace_back(i)  
...  
>>> len(v)  
10  
>>> for m in v:  
...     print(m.fData, end=' ')  
...  
0 1 2 3 4 5 6 7 8 9  
>>>
```



BERKELEY LAB

Runtime Cross-Inheritance

- **Cling's JIT compiles generated trampolines**
 - All proper C++ base classes can be inherited from
 - No need to select a subset of likely base classes
 - Only trampoline methods actually overridden
 - No overhead added to bound base class
 - Memory managed (copy, move, assign, destruct)

Runtime Cross-Inheritance

```
struct MyClass {  
    MyClass(int i) : fData(i) {}  
    virtual ~MyClass() {}  
    virtual int add(int i) {  
        return fData + i;  
    }  
    int fData;  
};
```

```
>>> import cppyy.gbl as CC  
>>> class PyMyClass(CC.MyClass):  
...     def add(self, i):  
...         return self.fData + 2*i  
...  
>>> m = PyMyClass(1)  
>>> CC.callb(m, 2)  
5  
>>>
```

The obvious next step ...

- Cross-inheritance allows Python classes in C++
 - Uniquely identifiable, memory managed
- C++ classes can be used as template argument
- Emergent property: *Python classes in templates!*

Thus obvious next step ...

```
struct MyClass {  
    MyClass(int i) : fData(i) {}  
    virtual ~MyClass() {}  
    virtual int add(int i) {  
        return fData + i;  
    }  
    int fData;  
};
```

```
>>> import cppyy.gbl as CC  
>>> class PyMyClass(CC.MyClass):  
...     def __init__(self, d, extra):  
...         super(PyMyClass, self).__init__(d)  
...         self.extra = extra  
...     def add(self, i):  
...         return self.fData + \  
...                 self.extra + 2*i  
...  
>>> v = \  
...     CC.std.vector[PyMyClass] ()  
...  
>>> v.push_back(PyMyClass(4, 42))  
>>> v.back().add(17)  
80  
>>>
```

Runtime Automatic Fallbacks

- Cling instantiates templates at runtime
- But Python types do no map uniquely, example:

Python	C++
type int	int8_t, uint8_t, short, unsigned short, int, unsigned int, long, unsigned long, long long, unsigned long long, int64_t, uint64_t, ...

- Solution: automatically fallback as needed

Runtime Automatic Fallbacks

```
template<typename T>
T passT(T t) {
    return t;
}
```

```
>>> import cppyy.gbl as CC
>>> type(1)
<class 'int'>
>>> CC.passT(1)
1
>>> CC.passT.__doc__
'int ::passT(int t)'
>>> type(2**64-1)
<class 'int'>
>>> CC.passT(2**64-1)
18446744073709551615
>>> CC.passT.__doc__
'unsigned long long ::passT(
    unsigned long long t)'
>>>
```



BERKELEY LAB

Runtime Callbacks

- **Cling's JIT compiles generated wrappers**
 - Type checked and memory managed
 - Errors (exceptions) can trace through both Python and C++
 - Python manages lifetime, C++ manages resources
 - Note: manage manually if C++ stores the function ptr
 - Supports C++ function pointers and std::function
- **Python can pass any callable**
 - Functions, lambda's, objects implementing __call__
 - Bound C++ functions and methods (w/o wrapper)

Runtime Callbacks

```
typedef int (*P)(int);

int callPtr(P f, int i) {
    return f(i);
}

typedef std::function<int(int)> F;

int callFun(const F& f, int i) {
    return f(i);
}
```

```
>>> import cppyy.gbl as CC
>>> def f(val):
...     return 2*val
...
>>> CC.callPtr(f, 2)
4
>>> CC.callFun(f, 3)
6
>>> CC.callPtr(lambda i: 5*i, 4)
20
>>> CC.callFun(lambda i: 6*i, 4)
24
>>>
```

Runtime Templated Callbacks

- **Modern Python3 supports “annotations”**
 - Very commonly used in any modern Python project
 - Type information used by IDEs, static checkers, etc.
 - Unused by (and mostly irrelevant to) the interpreter
 - Dictionary of (strings of) argument and return types
 - Strings are necessary for compound C++ types
- **Annotated functions can instantiate templates**
 - Incl. bound C++ functions (by definition “annotated”)

Runtime Templatized Callbacks

```
template<typename R,  
        typename... U,  
        typename... A>  
R callT(R(*f)(U...), A&&... a) {  
    return f(a...);  
}
```

```
>>> import cppyy.gbl as CC  
>>> def f(a: 'int') -> 'double':  
...     return 3.1415*a  
...  
>>> CC.callT(f, 2)  
6.283  
>>> def f(a: 'int', b: 'int') \  
-> 'int':  
...     return 3*a*b  
...  
>>> CC.callT(f, 6, 7)  
126  
>>>
```

Runtime Auto-downcast and Object Identity

- Always cast to the most derived C++ type
 - Involves a fake base and retrieving C++ RTTI
 - Custom RTTI implementation on MS Windows 64b
- Preserve identity Python proxy \Leftrightarrow C++ instance
 - Eases resource management / prevents dangling ptrs
 - Guarantees equal hashes for dictionary lookups
 - Alternative: specialize std::hash or __hash__
 - Enables cctor and assignment of cross-derived classes

Runtime Auto-downcast and Object Identity

```
struct Base {  
    virtual ~Base() {}  
};  
  
struct Derived : public Base {};  
  
Base* passB(Base* b) {  
    return b;  
}
```

```
>>> import cppyy.gbl as CC  
>>> d = CC.Derived()  
>>> b = CC.passB(d)  
>>> type(b) == CC.Derived  
True  
>>> d is b  
True  
>>>
```

Runtime Exceptions

- **Map exceptions derived from std::exception**
 - Python exceptions are not object instances
 - Python exception classes do not match C++ ones
- **Preserves C++ exception types**
 - Allows crossing multiple language layers
 - Provides trace showing full call stack

Note: can't mix compiled & JITed on all platforms

- To be fixed in a future version of Clang?



Runtime Exceptions

```
class MyException :  
    public std::exception {  
public:  
    const char* what() const throw() {  
        return "C++ failed";  
    }  
};  
  
void throw_error() {  
    throw MyException{};  
}
```

```
>>> import cppyy.gbl as CC  
>>> try:  
...     CC.throw_error()  
... except CC.MyException as e:  
...     print(e)  
...  
void ::throw_error() =>  
    MyException: C++ failed  
>>>
```

Runtime Unicode

- Python unicode encapsulates code points + codec
 - Defaults to UTF-8 (with BOM check)
- C++ is just all over the place, for example:
 - Byte-encoded w/o codec (e.g. std::string)
 - Code points w/o codec (e.g. std::u16string)
 - Wide char types w/ assumed codec (platform-specific)
- Python's type str != C++'s std::string
 - Developers still want interchangeable use by default ...

Runtime Unicode

```
template<class T>
std::string to_str(const T& chars) {
    char buf[12]; int n = 0;
    for (auto c : chars)
        buf[n++] = char(c);
    return std::string(buf, n-1);
}

std::string utf8_chinese() {
    auto chars = {0xe4, 0xb8, 0xad,
                  0xe6, 0x96, 0x87, 0};
    return to_str(chars);
}

std::string gbk_chinese() {
    auto chars = {0xd6, 0xd0, 0xce,
                  0xc4, 0};
    return to_str(chars);
}
```

```
>>> import cppyy.gbl as CC
>>> CC.utf8_chinese()
中文
>>> CC.gbk_chinese()
b'\xd6\xd0\xce\xc4'
>>> CC.gbk_chinese().decode('gbk')
中文
>>>
```

And much more, not directly runtime ...

- Classes, functions, (static) methods, operators, iterators, enums, single/multiple inheritance, shared/unique_ptr, STL pythonizations, ...
- Low-level C support (memory, arrays, ptr math, ...)
- Debug support (e.g. segfault -> Python exception)
- Customize with pythonizations, “freeze” binary distributions, cmake fragments for projects, ...
- See: <https://cppyy.readthedocs.io/en/latest/>

Yes, okay, runtime is great ... but
what about performance?

Performance Compared to Static Approaches

- No fundamental CPU performance difference

Note carefully that *everything* in Python is runtime: compile-time just means that the bindings *recipe* is compiled, not the actual bindings themselves!

- But heavy Cling/LLVM dependency:
 - ~25MB download cost; ~100MB memory overhead
 - Complex installation (and worse build)

Basic Performance Tests

```
void empty_call() {}

class Overload {
public:
    double add(int a, int b);
    double add(short a);
    double add(long a);
    double add(int a, int b, int c);
    double add(double a);
    double add(float a);
    double add(int a);
};

// benchmark example:
Overload obj;
for (size_t i=0; i < N; ++i)
    obj.add((double)i);
```

System:

Ubuntu 20.04.2 LTS
AMD EPYC 7702P 64-Core CPU
1TB of RAM

Setup:

gcc	9.3.0	(system)
pytest	6.2.4	(pypi)
benchmark:	3.4.1	(pypi)

Comparison:

cppyy	2.1.0	(pypi)
pybind11	2.7.1	(pypi)
swig	4.0.1	(system)
pypy-c	3.7.1	(system)



BERKELEY LAB

Basic Performance Test: empty call

Tool	Execution time (ns/call)*
C++ (Cling w/ -O2; out-of-line)	1.5
cppyy / pypy-c	16
swig (builtin)	27
cppyy / CPython	68
pybind11	68
swig (default)	104

- ⇒ Empty global function call is a pure overhead measure (zero work)
- ⇒ pypy-c slower than C++ b/c of global interpreter lock (GIL) release
- ⇒ "Builtin" swig trades functionality for speed
- ⇒ There is no obvious benefit to "static" over runtime bindings

(*) lower is better

Basic Performance Test: overload

Tool	Execution time (ms/call)*
C++ (Cling w/ -O2; out-of-line)	1.8E-6
cppyy / pypy-c	0.50
cppyy / CPython	1.25
swig (builtin)	1.29
swig (default)	4.23
pybind11	6.97

- ⇒ C++ overload is resolved at compile time, not based on dynamic type
- ⇒ Largest overhead: Python instance type checking (avoidable, but clumsy)
- ⇒ **There is no obvious benefit to “static” over runtime bindings**


(*) lower is better

Implementation: Bird's Eye View

Implementation

- Python offers hooks for C++ entity lookups, e.g.:
 - meta-classes for class creation
 - `__getattr__` for resolving attributes
 - `__getitem__`/`__call__` for template instantiations
- The hooks call into Cling for name lookup
 - All initial lookups are always *string-based*
- Access provided by address or through wrappers
 - Wrappers are C++ code to easily cover esoteric uses
 - Generalized interfaces to simplify downstream code

Implementation Flow: Lookup Example



Templates additionally use either `__getitem__` (explicit instantiation) or `__call__` (implicit).
Proxy creation calls into Cling for reflection information; wrappers are created on first use.

Implementation Flow: Call Example

```
m = MyClass()  
MyClass.Calc(42)
```

“Calc”
is unknown
↓
raise
Exception

“Calc” is
a template

call
create explicit
template

Calc<...>
is known
↓
Calc is
a function

generate
wrapper
→ call
Calc
↓
return
result

Wrappers of generated (and JITted) C++ are used to easily cover a range of C++isms, such as linkage of inline functions, overloaded operator new, default arguments, operator lookup, etc., etc. For simple cases in PyPy, direct FFI is used, for improved performance.

Conclusions

Current Limitations

- Complex and heavy Cling/LLVM dependency
- All C++ code enters single, global, translation unit
 - Significant slowdown for templates (Eigen, PCL, ...)
- PyPy/cppyy is significantly behind CPython/cppyy
- No MS Windows port for conda
- Significant Clang JIT limitations on MS Windows


Current Focus

- Add GPU (`cuda`) support to `cppyy`
- Bring PyPy / `cppyy` on par with CPython / `cppyy`
- Simplify installation / distribution

Conclusion

Runtime Python-C++ bindings are much more functional than similar static approaches
(and without loss in performance*)

(*) memory overhead is higher



That's all Folks!